

```
!pip install emoji
```

```
➦ Collecting emoji  
  Downloading emoji-2.14.1-py3-none-any.whl.metadata (5.7 kB)  
  Downloading emoji-2.14.1-py3-none-any.whl (590 kB)  
      ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 590.6/590.6 kB 17.6 MB/s eta 0:  
Installing collected packages: emoji  
Successfully installed emoji-2.14.1
```

```
import random
import numpy as np
import tensorflow as tf
import os
from tensorflow.keras.datasets import imdb
import pandas as pd
import matplotlib.pyplot as plt
import re
import emoji
from tqdm import tqdm
from sklearn.metrics import classification_report
import nltk
from nltk import pos_tag
from nltk.corpus import wordnet
from nltk.tokenize import TreebankWordTokenizer
from nltk.stem import WordNetLemmatizer
from tqdm import tqdm
from sklearn.model_selection import train_test_split
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, SimpleRNN, Dense, Dropout, Leaky
from tensorflow.keras.callbacks import EarlyStopping
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_s
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense, Dropout, LeakyReLU
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.layers import GRU
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import optuna
from tensorflow.keras.optimizers import Adam
```

SEED = 42

```
nltk.download('punkt')
nltk.download('wordnet')
nltk.download('averaged_perceptron_tagger_eng')
random.seed(SEED)
np.random.seed(SEED)
tf.random.set_seed(SEED)
os.environ['PYTHONHASHSEED'] = str(SEED)
os.environ['TF_DETERMINISTIC_OPS'] = '1'
```

```
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=10000)




word_index = imdb.get_word_index()
reverse_word_index = {value: key for key, value in word_index.items()}

def decode_review(encoded_review):
    return ' '.join([reverse_word_index.get(i - 3, '?') for i in encoded_review])

train_reviews = [decode_review(x) for x in x_train]
test_reviews = [decode_review(x) for x in x_test]

df = pd.DataFrame({
    "review": train_reviews + test_reviews,
    "label": list(y_train) + list(y_test)
})

df.head()
```

 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/17464789/17464789>  1s 0us/step  
 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/1641221/1641221>  1s 0us/step

	review	label
0	? this film was just brilliant casting locatio...	1
1	? big hair big boobs bad music and a giant saf...	0
2	? this has to be one of the worst films of the...	0
3	? the ? ? at storytelling the traditional sort...	1
4	? worst mistake of my life br br i picked this...	0

## ✓ 1-Text EDA

```
df.head()
```



	review	label
0	? this film was just brilliant casting locatio...	1
1	? big hair big boobs bad music and a giant saf...	0
2	? this has to be one of the worst films of the...	0
3	? the ? ? at storytelling the traditional sort...	1
4	? worst mistake of my life br br i picked this...	0

```
df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    review  50000 non-null    object
1    label   50000 non-null    int64
dtypes: int64(1), object(1)
memory usage: 781.4+ KB
```

```
df['label'].value_counts()
```



	count
label	
1	25000
0	25000

```
dtype: int64
```

```
df['review_length'] = df['review'].apply(lambda x: len(x.split()))
df['review_length'].describe()
```



review_length	
count	50000.000000
mean	234.755400
std	172.907439
min	7.000000
25%	129.000000
50%	176.000000
75%	285.000000
max	2494.000000

**dtype:** float64

```
df.isnull().sum()
```



	0
review	0
label	0
review_length	0

**dtype:** int64

```
df.sample(3)['review'].values
```



```
array(["? everyone should totally see this movie it's freaking scary but
doesn't resort to lame jump out at you just to surprise you and pass it
off as scary things it really is great see this freaking awesome movie the
director is stanley kubrick easily the greatest director who ever lived
every single one of his movies are masterpieces including this one the
shining is about this family that goes to a hotel in the ? ? as ? for the
winters and get ? in well the house is haunted the kid is psychic the
husband is easily ? by evil haunted ? and well hilarity ensues not really
it becomes this gripping thriller where stuff gets thrown at the viewer
from all different directions and it gets scary not just the classic
here's johnny scene it's memorable but can't speak for the whole movies
it's one of those things where words don't explain it adequately and you
just gotta see it so go on netflix and get it ? ?",
      "? the prey has an interesting history unless you remember the ads
```

for it in ? in june of 1984 you might have caught it on the movie channel back in summer 85 but little else is remembered the plot is your basic killer in the woods again but ironically this was filmed before friday the 13th the prey was actually shot sometime in 1978 according to one of the actors in an interview years later but released for about a week at some drive ins yes jim drive in showed this in june of ? but it has a dated look to it maybe they released it so later on to cash in on all the other terror films the market was ? with by 1984 now on the story it has some kind of back story a forest fire back in the 1940's leaves a lot of ? burned to death but one of their children survive our monster so flash forward to present day which would be 1978 we have an older middle age couple camping only to be ? by the monster the tag line for this picture claims its not human and its got an axe but an axe was only used in these first two killings now we have a bunch of teenagers who look like they in their mid ? camping we all know they are the prey and the monster knocks them of one by one for an 80 minute movie it seems longer we also have a lot of ? footage to fill in ? for the 80 mins overall for being out into an 80 's horror movie it looks way more 70's than ever hey the prey had potential to be a good horror killer in the woods movie but falls a little short it does however feature a pretty scary cool looking monster at the end and we have to wait till the last 2 minutes to see him side note the monster has gone on to star in the ? family movies in the 1990's",

"? he pulled the guys guts out his butt that's a spoof right no one really writes that it just happens like ? gone horribly wrong i think any way this movie must be a spoof because who would say they wrote that script otherwise can anyone imagine the entire cast sitting around as the director and writers go over the ? br br director says next our ? villain uses his 24 inch ? to ? our token creepy neighbor get this he is going to pull the guts out his ? br br brilliant the entire cast ? br br no way can that happen nobody writes that stupid gotta be a spoof br br i loved the part where the skinny ? gal beats the ? freak to death with the cast iron ? she finds on the floor of the cave i wasn't sure the ? cannibal types bothered to cook much maybe that explains why the ? was lying on the floor in the dark at just the right time to kill the ? hulk seems ironic that after the freaky guy had ? martial arts expert porn queens and a couple out doors type ? he falls so easily to the ? pan of a skinny ? girl next door br br what the heck is that richard ? guy doing in this did he fire his agent or something br br can anyone explain the ending to me please because i didn't get it either i can't quite figure why the nice hero girl wanted to kill the funny lady who was making her some tea never mind i don't want to know"],  
dtype=object)

```
df.shape
```

```
➡ (50000, 3)
```

```
df.columns
```

```
➡ Index(['review', 'label', 'review_length'], dtype='object')
```

```

word_index = imdb.get_word_index()
reverse_word_index = {value: key for key, value in word_index.items()}

def decode_review(encoded_review):
    return ' '.join([reverse_word_index.get(i - 3, '?') for i in encoded_review])

train_reviews = [decode_review(x) for x in x_train]
test_reviews = [decode_review(x) for x in x_test]

df = pd.DataFrame({
    "review": train_reviews + test_reviews,
    "label": list(y_train) + list(y_test)
})

label_map = {0: "Negative", 1: "Positive"}
df['label_name'] = df['label'].map(label_map)
sentiment_counts = df['label_name'].value_counts()

print(" Sentiment Distribution:")
print(sentiment_counts)

df['review_length'] = df['review'].apply(lambda x: len(x.split()))
length_stats = df['review_length'].describe()

print("\n Review Length Statistics (in words):")
print(length_stats)

```



Sentiment Distribution:

label\_name

Positive 25000

Negative 25000

Name: count, dtype: int64

Review Length Statistics (in words):

count 50000.000000

mean 234.755400

std 172.907439

min 7.000000

25% 129.000000

50% 176.000000

75% 285.000000

max 2494.000000

Name: review\_length, dtype: float64

## ✓ Best Practice to remove the text les then 10

```
df = df[df['review_length'] >= 10]
```

```
df.shape
```

```
↔ (49997, 4)
```

## 2-Text Preprocessing

### ✓ cleaning

```
url_pattern = re.compile(r"http\S+|www\S+|https\S+")
mention_pattern = re.compile(r'@\w+')
hashtag_pattern = re.compile(r'#\w+')
brackets_pattern = re.compile(r'\\(\\[\\{\\<|\\^()\\[\\]\\{\\}\\<>]*\\(\\)\\[\\]\\>|')
html_pattern = re.compile(r'<.*?>')
special_chars_pattern = re.compile(r'^a-zA-Z\\s')
digits_pattern = re.compile(r'\\d+')
repetition_pattern = re.compile(r'(.)\\1{2,}', re.DOTALL)
whitespace_pattern = re.compile(r'\\s+')
```

```
emoji_word_map = {
    "smiling_face_with_heart_eyes": "love",
    "face_with_tears_of_joy": "funny",
    "red_heart": "love",
    "crying_face": "sad",
    "face_with_symbols_on_mouth": "angry",
    "fire": "hot",
    "clapping_hands": "applause",
    "thumbs_up": "approve",
    "thumbs_down": "disapprove",
    "star_struck": "amazed",
    "grinning_face": "happy",
    "thinking_face": "thinking",
    "broken_heart": "heartbroken",
    "hundred_points": "perfect"
}
```

```
def replace_emoji_with_words(text):
```



```
text = emoji.demojize(text, delimiters=(" ", " "))
for code, word in emoji_word_map.items():
    text = text.replace(code, word)
return text

def clean_base_text(text):
    text = str(text).lower()
    text = replace_emoji_with_words(text)
    text = url_pattern.sub(' ', text)
    text = mention_pattern.sub(' ', text)
    text = hashtag_pattern.sub(' ', text)
    text = brackets_pattern.sub(' ', text)
    text = html_pattern.sub(' ', text)
    text = special_chars_pattern.sub(' ', text)
    text = digits_pattern.sub(' ', text)
    text = whitespace_pattern.sub(' ', text).strip()
    return text



def remove_excessive_repetition(text):
    return repetition_pattern.sub(r'\1\1', text)

def remove_useless_short_words(text):
    words = text.split()
    return ' '.join(word for word in words if len(word) > 2)

def full_cleaning_pipeline(text):
    text = clean_base_text(text)
    text = remove_excessive_repetition(text)
    text = remove_useless_short_words(text)
    return text

tqdm.pandas()
df['cleaned_text'] = df['review'].progress_apply(full_cleaning_pipeline)

df[['review', 'cleaned_text']].sample(5)
```

 100% |  | 49997/49997 [00:50<00:00, 989.66it/s]

	review	cleaned_text
5480	? i don't know why i keep doing this to myself...	don know why keep doing this myself keep defen...
995	? i think this still is the best routine there...	think this still the best routine there are so...
44278	? i decided to hire out this movie along with ...	decided hire out this movie along with few oth...
44278	? i couldn't believe the eye candy from	couldn believe the eye candy from start

## ✓ Lemmatization

```

lemmatizer = WordNetLemmatizer()
tokenizer = TreebankWordTokenizer()

def get_wordnet_pos(tag):
    if tag.startswith('J'):
        return wordnet.ADJ
    elif tag.startswith('V'):
        return wordnet.VERB
    elif tag.startswith('N'):
        return wordnet.NOUN
    elif tag.startswith('R'):
        return wordnet.ADV
    else:
        return wordnet.NOUN

def conservative_lemmatize(text):
    tokens = tokenizer.tokenize(text)
    tagged_tokens = pos_tag(tokens)
    lemmatized_tokens = [
        lemmatizer.lemmatize(token, get_wordnet_pos(tag))
        for token, tag in tagged_tokens
    ]
    return ' '.join(lemmatized_tokens)

tqdm.pandas()
df['lemmatized_text'] = df['cleaned_text'].progress_apply(conservative_lemmatize)
df[['cleaned_text', 'lemmatized_text']].sample(5)

```

```

[ntlk_data] Downloading package punkt to /root/nltk_data...
[ntlk_data]   Unzipping tokenizers/punkt.zip.
[ntlk_data] Downloading package wordnet to /root/nltk_data...
[ntlk_data] Downloading package averaged_perceptron_tagger_eng to
[ntlk_data]   /root/nltk_data...
[ntlk_data]   Unzipping taggers/averaged_perceptron_tagger_eng.zip.
100%|██████████| 49997/49997 [05:05<00:00, 163.46it/s]

```

	cleaned_text	lemmatized_text
8512	first this movie seems bad that almost fell th...	first this movie seem bad that almost fell the...
5097	the movie starts off setting where not surpris...	the movie start off set where not surprisingly...
28531	this thriller has many twists and turns had th...	this thriller have many twist and turn have th...
21000	know know plan from outer space the	know know plan from out space the bad

## ✓ Text Vectorization – Tokenization and Padding

[illegible]

## ✓ Build and Train Simple RNN Model

```
model_rnn = Sequential([
    Embedding(input_dim=10000, output_dim=64),

    Bidirectional(SimpleRNN(units=64, return_sequences=True)),
    Dropout(0.3),

    Bidirectional(SimpleRNN(units=32)),
    Dropout(0.3),















    Dense(32),
    LeakyReLU(alpha=0.01),

    Dense(1, activation='sigmoid')
])

model_rnn.compile(
    loss='binary_crossentropy',
    optimizer='adam',
    metrics=['accuracy']
)

history_rnn = model_rnn.fit(
    X_train_pad, y_train,
    validation_data=(X_test_pad, y_test),
    epochs=60,
    batch_size=128,
    callbacks=[EarlyStopping(patience=5, restore_best_weights=True)],
    verbose=1
)
```

```

Epoch 1/60
/usr/local/lib/python3.11/dist-packages/keras/src/layers/activations/leaky_
warnings.warn(
313/313  193s 594ms/step - accuracy: 0.5175 - loss: 0.6
Epoch 2/60
313/313  193s 565ms/step - accuracy: 0.7435 - loss: 0.5
Epoch 3/60
313/313  195s 542ms/step - accuracy: 0.7161 - loss: 0.5
Epoch 4/60
313/313  173s 553ms/step - accuracy: 0.7779 - loss: 0.4
Epoch 5/60
313/313  214s 593ms/step - accuracy: 0.6188 - loss: 0.6
Epoch 6/60
313/313  214s 632ms/step - accuracy: 0.6730 - loss: 0.5
Epoch 7/60
313/313  220s 688ms/step - accuracy: 0.7350 - loss: 0.5
Epoch 8/60
313/313  172s 548ms/step - accuracy: 0.8301 - loss: 0.4
Epoch 9/60
313/313  203s 552ms/step - accuracy: 0.8580 - loss: 0.3
Epoch 10/60
313/313  203s 555ms/step - accuracy: 0.8498 - loss: 0.3
Epoch 11/60
313/313  199s 546ms/step - accuracy: 0.7366 - loss: 0.5
Epoch 12/60
313/313  199s 538ms/step - accuracy: 0.7349 - loss: 0.5
Epoch 13/60
313/313  171s 547ms/step - accuracy: 0.8355 - loss: 0.4
Epoch 14/60
313/313  199s 539ms/step - accuracy: 0.8318 - loss: 0.3

```

```

y_pred_prob = model_rnn.predict(X_test_pad).flatten()
y_pred = (y_pred_prob > 0.5).astype("int32")

```

```

313/313  27s 85ms/step

```

```

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print(f"Accuracy:  {accuracy:.4f}")
print(f" Precision: {precision:.4f}")
print(f" Recall:    {recall:.4f}")
print(f" F1 Score:   {f1:.4f}")

```

```

➦ Accuracy:  0.8261
    Precision: 0.8227
    Recall:    0.8314
    F1 Score:  0.8270

```


## ✓ Accuracy / Loss Curves

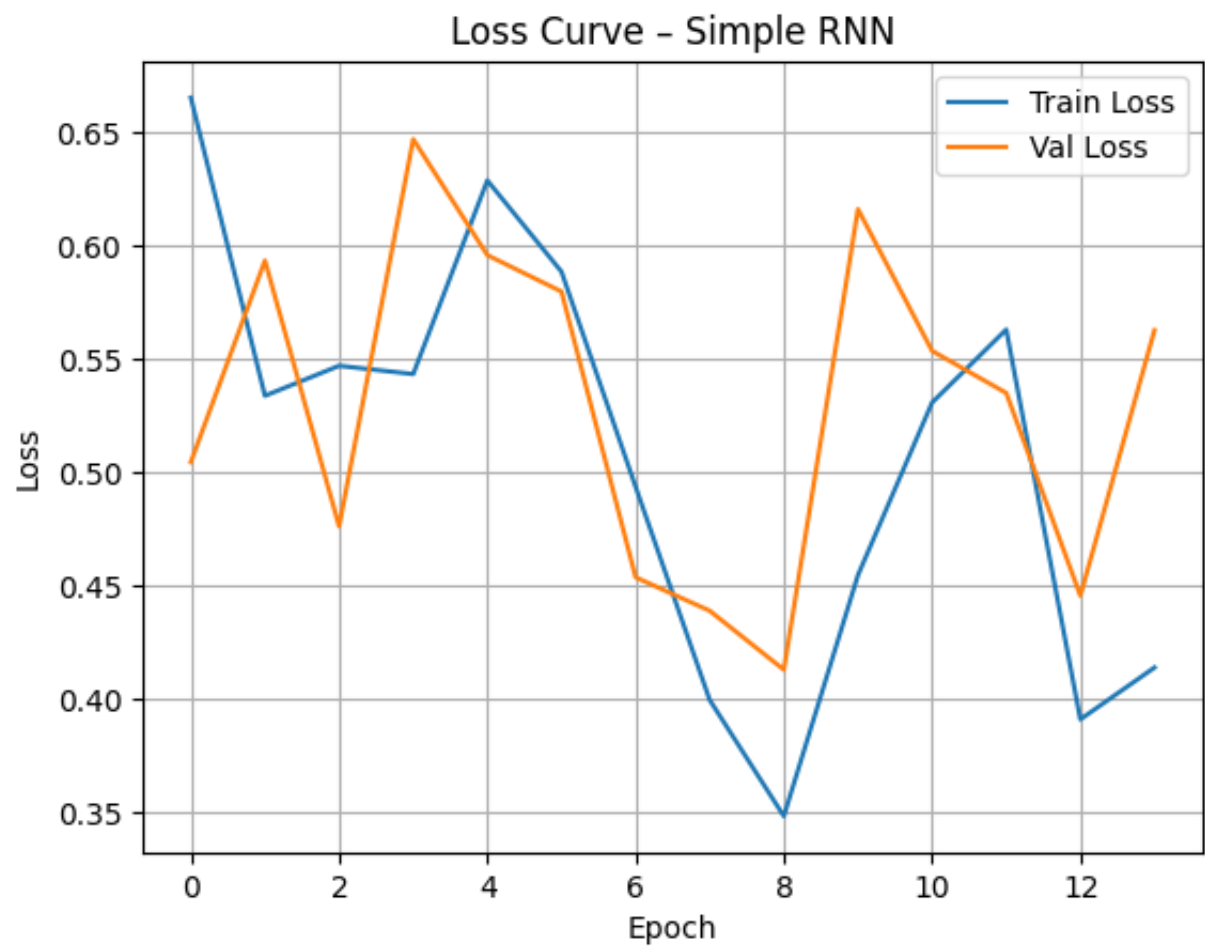

```

plt.plot(history_rnn.history['accuracy'], label='Train Accuracy')
plt.plot(history_rnn.history['val_accuracy'], label='Val Accuracy')
plt.title('Accuracy Curve – Simple RNN')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)
plt.show()

plt.plot(history_rnn.history['loss'], label='Train Loss')
plt.plot(history_rnn.history['val_loss'], label='Val Loss')
plt.title('Loss Curve – Simple RNN')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.show()

```






## ✓ Confusion Matrix



```
y_pred_prob = model_rnn.predict(X_test_pad)
y_pred = (y_pred_prob > 0.5).astype("int32")

cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Negative",
disp.plot(cmap='Blues')
plt.title("Confusion Matrix - Simple RNN")
plt.show()
```

313/313 — 22s 72ms/step



## ✓ Classification Report

```
print(classification_report(y_test, y_pred, target_names=["Negative", "Positive"])
```

	precision	recall	f1-score	support
Negative	0.83	0.82	0.83	5000
Positive	0.82	0.83	0.83	5000
accuracy			0.83	10000
macro avg	0.83	0.83	0.83	10000
weighted avg	0.83	0.83	0.83	10000

## ✓ Qualitative Analysis

High-Confidence Correct Predictions

```

y_pred_prob = model_rnn.predict(X_test_pad).flatten()
y_pred = (y_pred_prob > 0.5).astype("int32")

correct_mask = (y_pred == y_test)
correct_probs = y_pred_prob[correct_mask]
correct_texts = X_test_texts[correct_mask]
correct_labels = y_test[correct_mask]

top_correct_indices = np.argsort(correct_probs)[-5:] [::-1]

print(" Top High-Confidence Correct Predictions:\n")
for i in top_correct_indices:
    print(" Review:")
    print(correct_texts[i][:500])
    print(f"Actual Label: {correct_labels[i]}, Predicted Prob: {correct_probs[i]}")
    print("-" * 60)

```

 **313/313**  **20s** 63ms/step

Top High-Confidence Correct Predictions:

Review:

excellent blend music light comedy and drama with picture perfect performan  
Actual Label: 1, Predicted Prob: 0.9768

Review:

excellent comedy star dudley moore support minnelli and good speaking john  
Actual Label: 1, Predicted Prob: 0.9768

Review:

excellent story telling and cinematography poignant bite social commentary  
Actual Label: 1, Predicted Prob: 0.9767

Review:

deathtrap give you twist every turn every single turn fact it big problem t  
Actual Label: 1, Predicted Prob: 0.9767

Review:

refresh breath air when movie actually give you story line with begin middl  
Actual Label: 1, Predicted Prob: 0.9766

## low-Confidence Predictions

```
wrong_preds = (y_test != y_pred.flatten())
wrong_reviews = X_test_texts[wrong_preds]

for i in range(3):
    print(" Review:")
    print(wrong_reviews[i])
    print(f"Actual: {y_test[wrong_preds][i]}, Predicted: {y_pred[wrong_preds][i]}")
    print("-" * 60)
```

```
➞ Review:
have be pleasantly surprised sandra performance miss decide give murder num
Actual: 0, Predicted: 1
-----
Review:
have never hear larry before but judge this effort into write and direct sh
Actual: 0, Predicted: 1
-----
Review:
two star amanda plummer look like young version her father christopher plum
Actual: 0, Predicted: 1
-----
```

## ✓ Long Short-Term Memory -LSTM-

```
model_lstm = Sequential([
    Embedding(input_dim=10000, output_dim=64, input_length=300),
    LSTM(units=64, return_sequences=False),
    Dropout(0.4),
    Dense(32),
    LeakyReLU(alpha=0.01),
    Dense(1, activation='sigmoid')
])

model_lstm.compile(
    loss='binary_crossentropy',
    optimizer='adam',
    metrics=['accuracy']
)
```

```
history_lstm = model_lstm.fit(
    X_train_pad, y_train,
    validation_data=(X_test_pad, y_test),
    epochs=60,
    batch_size=128,
    callbacks=[EarlyStopping(patience=5, restore_best_weights=True)],
    verbose=1
)
```

```
↔ Epoch 1/60
313/313 ————— 179s 559ms/step - accuracy: 0.4989 - loss: 0.6
Epoch 2/60
313/313 ————— 148s 472ms/step - accuracy: 0.5175 - loss: 0.6
Epoch 3/60
313/313 ————— 213s 507ms/step - accuracy: 0.5328 - loss: 0.6
Epoch 4/60
313/313 ————— 159s 508ms/step - accuracy: 0.5523 - loss: 0.6
Epoch 5/60
313/313 ————— 158s 505ms/step - accuracy: 0.5540 - loss: 0.6
Epoch 6/60
313/313 ————— 158s 506ms/step - accuracy: 0.5516 - loss: 0.6
Epoch 7/60
313/313 ————— 206s 517ms/step - accuracy: 0.5640 - loss: 0.6
Epoch 8/60
313/313 ————— 152s 486ms/step - accuracy: 0.7440 - loss: 0.5
Epoch 9/60
313/313 ————— 201s 483ms/step - accuracy: 0.8918 - loss: 0.2
Epoch 10/60
313/313 ————— 151s 483ms/step - accuracy: 0.9306 - loss: 0.1
Epoch 11/60
313/313 ————— 159s 508ms/step - accuracy: 0.9506 - loss: 0.1
Epoch 12/60
313/313 ————— 202s 508ms/step - accuracy: 0.9616 - loss: 0.1
Epoch 13/60
313/313 ————— 194s 484ms/step - accuracy: 0.9747 - loss: 0.0
Epoch 14/60
313/313 ————— 159s 509ms/step - accuracy: 0.9809 - loss: 0.0
```

## ✓ Evaluation of LSTM Model

```
y_pred_prob = model_lstm.predict(X_test_pad).flatten()
y_pred = (y_pred_prob > 0.5).astype("int32")

print(f" Accuracy:  {accuracy_score(y_test, y_pred):.4f}")
print(f" Precision: {precision_score(y_test, y_pred):.4f}")
print(f" Recall:    {recall_score(y_test, y_pred):.4f}")
print(f" F1 Score:   {f1_score(y_test, y_pred):.4f}")
```

↗ **313/313** ————— **24s** 74ms/step

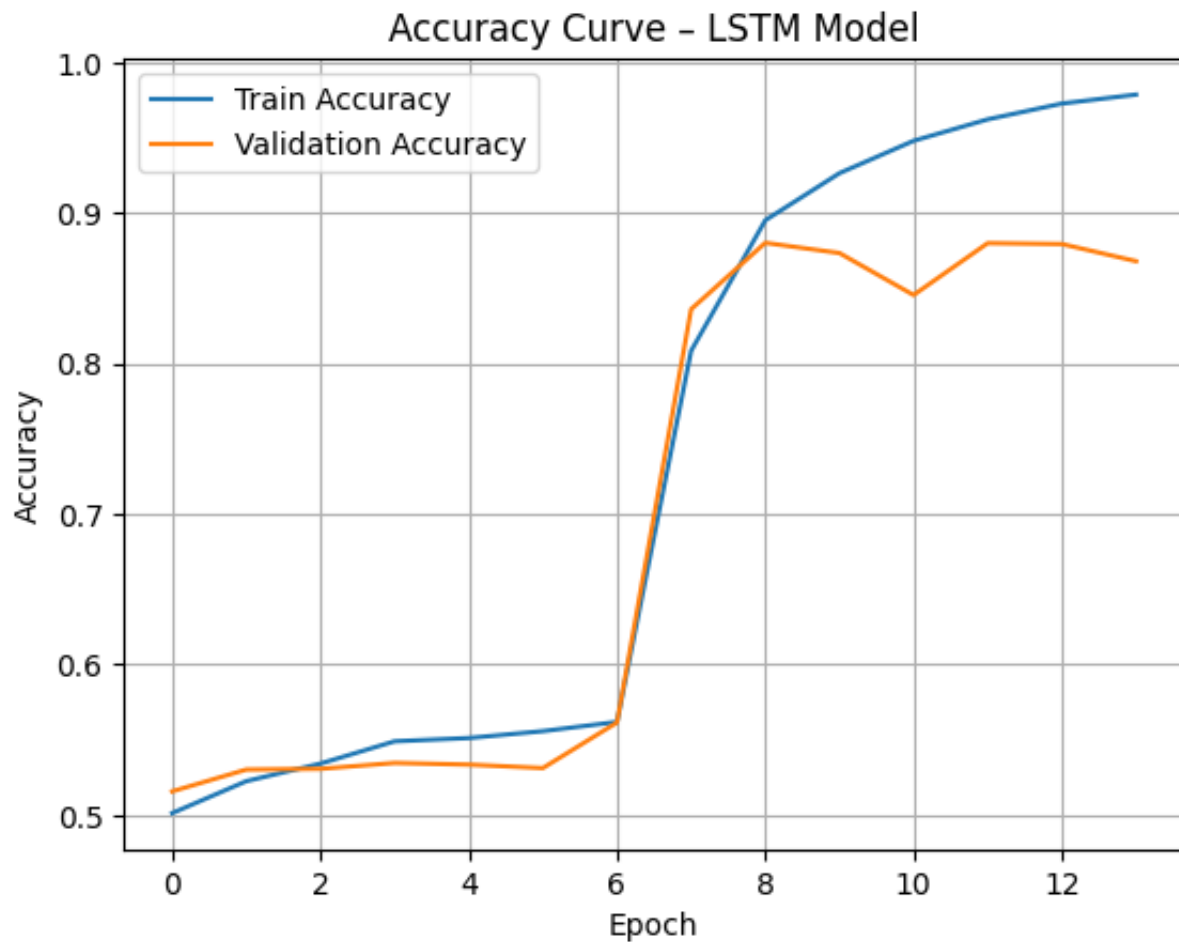
Accuracy: 0.8804  
Precision: 0.8737  
Recall: 0.8894  
F1 Score: 0.8815

**313/313** ————— **12s** 37ms/step

Accuracy: 0.8804  
Precision: 0.8737  
Recall: 0.8894  
F1 Score: 0.8815


Plot Training & Validation Accuracy

```
plt.plot(history_lstm.history['accuracy'], label='Train Accuracy')
plt.plot(history_lstm.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy Curve - LSTM Model')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)
plt.show()
```



## ✓ Plot Training & Validation Loss

```
plt.plot(history_lstm.history['loss'], label='Train Loss')
plt.plot(history_lstm.history['val_loss'], label='Validation Loss')
plt.title('Loss Curve - LSTM Model')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.show()
```




## ✓ Confusion Matrix



```
y_pred_prob = model_lstm.predict(X_test_pad).flatten()
y_pred = (y_pred_prob > 0.5).astype("int32")

cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Negative",
disp.plot(cmap='Blues')
plt.title("Confusion Matrix - LSTM Model")
plt.show()
```

313/313 — 11s 36ms/step



## ✓ Classification Report

```
print(classification_report(y_test, y_pred, target_names=["Negative", "Positive"])
```

```

precision    recall  f1-score   support

Negative     0.89     0.87     0.88     5000
Positive     0.87     0.89     0.88     5000

accuracy          0.88     10000
macro avg         0.88     0.88     0.88     10000
weighted avg      0.88     0.88     0.88     10000

```

## High-Confidence Wrong Predictions (Qualitative Error Analysis)

```

wrong_mask = (y_pred != y_test)
wrong_conf = y_pred_prob[wrong_mask]
wrong_texts = X_test_texts[wrong_mask]

top_wrong = np.argsort(wrong_conf)[-3:][::-1]

print(" High-Confidence Wrong Predictions:\n")
for i in top_wrong:
    print("Review:")
    print(wrong_texts[i][:500]) # Use array indexing instead of .iloc
    print(f"Actual: {y_test[wrong_mask][i]}, Predicted Prob: {wrong_conf[i]:.4f}")
    print("-" * 60)

```

High-Confidence Wrong Predictions:

```

Review:
firmly believe that the best oscar ceremony recent year be for two reason h
Actual: 0, Predicted Prob: 1.0000
-----
Review:
the story go something like this small town girl katie jessica simpson deci
Actual: 0, Predicted Prob: 1.0000
-----
Review:
saw this movie just now not when be release and best picture the year here
Actual: 0, Predicted Prob: 1.0000
-----

```

## High-Confidence Correct Predictions

```
correct_mask = (y_pred == y_test)
correct_conf = y_pred_prob[correct_mask]
correct_texts = X_test_texts[correct_mask]

top_correct = np.argsort(correct_conf)[-3:][::-1]

print("High-Confidence Correct Predictions:\n")
for i in top_correct:
    print("Review:")
    print(correct_texts[i][:500]) # Access using array indexing
    print(f"Actual: {y_test[correct_mask][i]}, Predicted Prob: {correct_conf[i]}")
    print("-" * 60)
```

⇒ High-Confidence Correct Predictions:

Review:  
the notorious bettie page gretchen mol taylor chris bauer jar harris sarah  
Actual: 1, Predicted Prob: 1.0000  
-----

Review:  
one reason pixar have endure well and be successful that while their film r  
Actual: 1, Predicted Prob: 1.0000  
-----

Review:  
be see the year the matrix with the release two sequel and computer game th  
Actual: 1, Predicted Prob: 1.0000  
-----

## ✓ GRU Model

```
model_gru = Sequential([
    Embedding(input_dim=10000, output_dim=64, input_length=300),
    GRU(units=64, return_sequences=False),
    Dropout(0.4),
    Dense(32),
    LeakyReLU(alpha=0.01),
```

```

        Dense(1, activation='sigmoid')
    ])

model_gru.compile(
    loss='binary_crossentropy',
    optimizer='adam',
    metrics=['accuracy']
)

history_gru = model_gru.fit(
    X_train_pad, y_train,
    validation_data=(X_test_pad, y_test),
    epochs=60,
    batch_size=128,
    callbacks=[EarlyStopping(patience=5, restore_best_weights=True)],
    verbose=1
)

```

→ /usr/local/lib/python3.11/dist-packages/keras/src/layers/core/embedding.py:  
 warnings.warn(  
 /usr/local/lib/python3.11/dist-packages/keras/src/layers/activations/leaky\_  
 warnings.warn(  
 Epoch 1/60  
**313/313** ————— **234s** 723ms/step - accuracy: 0.5057 - loss: 0.6  
 Epoch 2/60  
**313/313** ————— **235s** 639ms/step - accuracy: 0.5080 - loss: 0.6  
 Epoch 3/60  
**313/313** ————— **209s** 668ms/step - accuracy: 0.5378 - loss: 0.6  
 Epoch 4/60  
**313/313** ————— **253s** 639ms/step - accuracy: 0.7692 - loss: 0.4  
 Epoch 5/60  
**313/313** ————— **205s** 654ms/step - accuracy: 0.9101 - loss: 0.2  
 Epoch 6/60  
**313/313** ————— **260s** 649ms/step - accuracy: 0.9380 - loss: 0.1  
 Epoch 7/60  
**313/313** ————— **203s** 648ms/step - accuracy: 0.9553 - loss: 0.1  
 Epoch 8/60  
**313/313** ————— **261s** 647ms/step - accuracy: 0.9681 - loss: 0.1  
 Epoch 9/60  
**313/313** ————— **265s** 659ms/step - accuracy: 0.9785 - loss: 0.0  
 Epoch 10/60  
**313/313** ————— **260s** 652ms/step - accuracy: 0.9836 - loss: 0.0

Generate Predictions

```
y_pred_prob = model_gru.predict(X_test_pad).flatten()
y_pred = (y_pred_prob > 0.5).astype("int32")
```


↔ 313/313 ————— 12s 39ms/step

```
print(f" Accuracy: {accuracy_score(y_test, y_pred):.4f}")
print(f" Precision: {precision_score(y_test, y_pred):.4f}")
print(f" Recall: {recall_score(y_test, y_pred):.4f}")
print(f" F1 Score: {f1_score(y_test, y_pred):.4f}")
```

↔ Accuracy: 0.8983  
Precision: 0.8955  
Recall: 0.9018  
F1 Score: 0.8987


Plot Confusion Matrix

```
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Negative",
disp.plot(cmap='Greens')
plt.title("Confusion Matrix - GRU Model")
plt.show()
```




Accuracy Curve

```
plt.plot(history_gru.history['accuracy'], label='Train Accuracy')
plt.plot(history_gru.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy Curve - GRU Model')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)
plt.show()
```



Loss Curve

```
plt.plot(history_gru.history['loss'], label='Train Loss')
plt.plot(history_gru.history['val_loss'], label='Validation Loss')
plt.title('Loss Curve - GRU Model')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.show()
```



## ✓ Compare All Models: Accuracy & Loss Curves

Accuracy Comparison Plot



```
plt.figure(figsize=(10,6))
plt.plot(history_rnn.history['val_accuracy'], label='RNN - Validation Accuracy')
plt.plot(history_lstm.history['val_accuracy'], label='LSTM - Validation Accuracy')
plt.plot(history_gru.history['val_accuracy'], label='GRU - Validation Accuracy')

plt.title(' Validation Accuracy Comparison')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)
plt.show()
```

 [Show hidden output](#)

## Loss Comparison Plot

```
plt.figure(figsize=(10,6))

plt.plot(history_rnn.history['val_loss'], label='RNN - Validation Loss')
plt.plot(history_lstm.history['val_loss'], label='LSTM - Validation Loss')
plt.plot(history_gru.history['val_loss'], label='GRU - Validation Loss')

plt.title(' Validation Loss Comparison')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.show()
```

## ✓ models save

```
model_gru.save("gru_model.h5")
model_lstm.save("lstm_model.h5")
model_rnn.save("rnn_model.h5")
```

## ✓ Hyperparameter Optimization

After evaluating all models, GRU achieved the highest validation performance with efficient training time. Therefore, it was selected for hyperparameter optimization using Optuna to further maximize its accuracy and generalization.

```
!pip install optuna
```

```

⇌ Collecting optuna
  Downloading optuna-4.3.0-py3-none-any.whl.metadata (17 kB)
Collecting alembic>=1.5.0 (from optuna)
  Downloading alembic-1.15.2-py3-none-any.whl.metadata (7.3 kB)
Collecting colorlog (from optuna)
  Downloading colorlog-6.9.0-py3-none-any.whl.metadata (10 kB)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-pack
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11
Requirement already satisfied: sqlalchemy>=1.4.2 in /usr/local/lib/python3.
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packa
Requirement already satisfied: PyYAML in /usr/local/lib/python3.11/dist-pac
Requirement already satisfied: Mako in /usr/lib/python3/dist-packages (from
Requirement already satisfied: typing-extensions>=4.12 in /usr/local/lib/py
Requirement already satisfied: greenlet>=1 in /usr/local/lib/python3.11/dis
Downloading optuna-4.3.0-py3-none-any.whl (386 kB)
  ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 386.6/386.6 kB 18.3 MB/s eta 0:
Downloading alembic-1.15.2-py3-none-any.whl (231 kB)
  ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 231.9/231.9 kB 21.6 MB/s eta 0:
Downloading colorlog-6.9.0-py3-none-any.whl (11 kB)
Installing collected packages: colorlog, alembic, optuna
Successfully installed alembic-1.15.2 colorlog-6.9.0 optuna-4.3.0

```

```
def objective(trial):
    model = Sequential()

    embedding_dim = trial.suggest_categorical('embedding_dim', [32, 64, 128])
    model.add(Embedding(input_dim=10000, output_dim=embedding_dim, input_length=

    gru_units = trial.suggest_int('gru_units', 32, 128, step=32)
    model.add(GRU(units=gru_units, return_sequences=False))

    dropout_rate = trial.suggest_float('dropout', 0.2, 0.5, step=0.1)
    model.add(Dropout(rate=dropout_rate))

    model.add(Dense(1, activation='sigmoid'))

    learning_rate = trial.suggest_float("lr", 1e-4, 1e-2, log=True)
    model.compile(
        loss='binary_crossentropy',
        optimizer=Adam(learning_rate=learning_rate),
        metrics=['accuracy']
    )

    batch_size = trial.suggest_categorical("batch_size", [64, 128, 256])
    history = model.fit(
        X_train_pad, y_train,
        validation_data=(X_test_pad, y_test),
        epochs=5,
        batch_size=batch_size,
        verbose=0
    )

    val_accuracy = history.history["val_accuracy"][-1]
    return val_accuracy
```

```
study = optuna.create_study(direction="maximize")
study.optimize(objective, n_trials=10)
```

```
[I 2025-04-18 13:34:34,230] A new study created in memory with name: no-nam
/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/embedding.py:
warnings.warn(
[I 2025-04-18 13:35:18,884] Trial 0 finished with value: 0.8898000121116638
[I 2025-04-18 13:35:40,368] Trial 1 finished with value: 0.890500009059906
[I 2025-04-18 13:36:22,636] Trial 2 finished with value: 0.6812999844551086
[I 2025-04-18 13:36:44,118] Trial 3 finished with value: 0.5289999842643738
[I 2025-04-18 13:36:59,823] Trial 4 finished with value: 0.5200999975204468
[I 2025-04-18 13:37:24,329] Trial 5 finished with value: 0.5073000192642212
[I 2025-04-18 13:38:03,015] Trial 6 finished with value: 0.8946999907493591
[I 2025-04-18 13:38:29,613] Trial 7 finished with value: 0.8883000016212463
[I 2025-04-18 13:39:08,501] Trial 8 finished with value: 0.8877000212669373
[I 2025-04-18 13:39:22,394] Trial 9 finished with value: 0.727400004863739
```

```
print(" Best Hyperparameters:")
for key, value in study.best_params.items():
    print(f"{key}: {value}")

print(f"\n Best Validation Accuracy: {study.best_value:.4f}")
```

```
[I 2025-04-18 13:39:22,394] Best Hyperparameters:
embedding_dim: 32
gru_units: 32
dropout: 0.5
lr: 0.002925248561355503
batch_size: 64

Best Validation Accuracy: 0.8947
```

Hyperparameter optimization was conducted using Optuna with a search space over embedding\_dim, gru\_units, dropout, batch\_size, and learning\_rate. The best configuration achieved a validation accuracy of {study.best\_value:.4f}.

After automated hyperparameter tuning using Optuna, we rebuilt the GRU model using the best configuration and trained it for 10 full epochs. This yielded the final optimized performance reported."

## ✓ Hyperparameter Optimization on GRU Model

```
model_gru_opt = Sequential([
    Embedding(input_dim=10000, output_dim=32, input_length=300),
    GRU(units=32, return_sequences=False),
    Dropout(0.5),
    Dense(32),
    LeakyReLU(alpha=0.01),
    Dense(1, activation='sigmoid')
])

optimizer = Adam(learning_rate=0.00293)

model_gru_opt.compile(
    loss='binary_crossentropy',
    optimizer=optimizer,
    metrics=['accuracy']
)

history_gru_opt = model_gru_opt.fit(
    X_train_pad, y_train,
    validation_data=(X_test_pad, y_test),
    epochs=60,
    batch_size=64,
    callbacks=[EarlyStopping(patience=10, restore_best_weights=True)],
    verbose=1
)
```

```

➡ Epoch 1/60
625/625 ————— 10s 13ms/step - accuracy: 0.4992 - loss: 0.693
Epoch 2/60
625/625 ————— 8s 13ms/step - accuracy: 0.5103 - loss: 0.6916
Epoch 3/60
625/625 ————— 8s 12ms/step - accuracy: 0.8040 - loss: 0.4379
Epoch 4/60
625/625 ————— 8s 13ms/step - accuracy: 0.9206 - loss: 0.2159
Epoch 5/60
625/625 ————— 8s 13ms/step - accuracy: 0.9465 - loss: 0.1572
Epoch 6/60
625/625 ————— 8s 13ms/step - accuracy: 0.9628 - loss: 0.1140
Epoch 7/60
625/625 ————— 8s 13ms/step - accuracy: 0.9705 - loss: 0.0952
Epoch 8/60
625/625 ————— 8s 13ms/step - accuracy: 0.9795 - loss: 0.0633
Epoch 9/60
625/625 ————— 8s 12ms/step - accuracy: 0.9864 - loss: 0.0465
Epoch 10/60
625/625 ————— 8s 13ms/step - accuracy: 0.9880 - loss: 0.0434
Epoch 11/60
625/625 ————— 8s 13ms/step - accuracy: 0.9901 - loss: 0.0339
Epoch 12/60
625/625 ————— 8s 12ms/step - accuracy: 0.9917 - loss: 0.0272
Epoch 13/60
625/625 ————— 8s 13ms/step - accuracy: 0.9946 - loss: 0.0199

```

```

y_pred_prob = model_gru_opt.predict(X_test_pad).flatten()
y_pred = (y_pred_prob > 0.5).astype("int32")

```

```

➡ 313/313 ————— 2s 5ms/step

```

```

print(f" Accuracy: {accuracy_score(y_test, y_pred):.4f}")
print(f" Precision: {precision_score(y_test, y_pred):.4f}")
print(f" Recall: {recall_score(y_test, y_pred):.4f}")
print(f" F1 Score: {f1_score(y_test, y_pred):.4f}")


```

```


➡ Accuracy: 0.8711
Precision: 0.8425
Recall: 0.9128
F1 Score: 0.8763

```

```
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Negative",
disp.plot(cmap='Greens')
plt.title("Confusion Matrix - GRU Model")
plt.show()
```




```
plt.plot(history_gru_opt.history['accuracy'], label='Train Accuracy')
plt.plot(history_gru_opt.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy Curve - GRU Model')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)
plt.show()
```





```
plt.plot(history_gru_opt.history['loss'], label='Train Loss')
plt.plot(history_gru_opt.history['val_loss'], label='Validation Loss')
plt.title('Loss Curve - GRU Model')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.show()
```



## ✓ Bidirectional GRU

```

model_bi_gru = Sequential([
    Embedding(input_dim=10000, output_dim=32, input_length=300),
    Bidirectional(GRU(units=32, return_sequences=False)),
    Dropout(0.5),
    Dense(32),
    LeakyReLU(alpha=0.01),
    Dense(1, activation='sigmoid')
])

optimizer = Adam(learning_rate=0.00293)

model_bi_gru.compile(
    loss='binary_crossentropy',
    optimizer=optimizer,
    metrics=['accuracy']
)

history_bi_gru = model_bi_gru.fit(
    X_train_pad, y_train,
    validation_data=(X_test_pad, y_test),
    epochs=60,
    batch_size=64,
    callbacks=[EarlyStopping(patience=5, restore_best_weights=True)],
    verbose=1
)

```

```

↔ Epoch 1/60
625/625 ————— 16s 21ms/step - accuracy: 0.6878 - loss: 0.565
Epoch 2/60
625/625 ————— 13s 21ms/step - accuracy: 0.8866 - loss: 0.282
Epoch 3/60
625/625 ————— 13s 21ms/step - accuracy: 0.9225 - loss: 0.207
Epoch 4/60
625/625 ————— 13s 21ms/step - accuracy: 0.9452 - loss: 0.153
Epoch 5/60
625/625 ————— 13s 21ms/step - accuracy: 0.9592 - loss: 0.118
Epoch 6/60
625/625 ————— 13s 21ms/step - accuracy: 0.9705 - loss: 0.089
Epoch 7/60
625/625 ————— 13s 21ms/step - accuracy: 0.9731 - loss: 0.080


```

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

y_pred_prob = model_bi_gru.predict(X_test_pad).flatten()
y_pred = (y_pred_prob > 0.5).astype("int32")
```

 **313/313** ————— **3s** 8ms/step

```
print(f" Accuracy: {accuracy_score(y_test, y_pred):.4f}")
print(f" Precision: {precision_score(y_test, y_pred):.4f}")
print(f" Recall: {recall_score(y_test, y_pred):.4f}")
print(f" F1 Score: {f1_score(y_test, y_pred):.4f}")
```

 Accuracy: 0.8954  
Precision: 0.8965  
Recall: 0.8940  
F1 Score: 0.8953



```
def plot_training_history(history, model_name="Model"):
    plt.figure(figsize=(14, 5))

    plt.subplot(1, 2, 1)
    plt.plot(history.history['accuracy'], label='Train Accuracy')
    plt.plot(history.history['val_accuracy'], label='Val Accuracy')
    plt.title(f'{model_name} Accuracy')
    plt.xlabel('Epoch')
    plt.ylabel('Accuracy')
    plt.legend()

    plt.subplot(1, 2, 2)
    plt.plot(history.history['loss'], label='Train Loss')
    plt.plot(history.history['val_loss'], label='Val Loss')
    plt.title(f'{model_name} Loss')
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.legend()

    plt.tight_layout()
    plt.show()

plot_training_history(history_bi_gru, model_name="Bidirectional GRU")
```



```

y_pred_prob = model_bi_gru.predict(X_test_pad)
y_pred = (y_pred_prob > 0.5).astype("int32")

print("Classification Report:")
print(classification_report(y_test, y_pred, digits=4))

cm = confusion_matrix(y_test, y_pred)


plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['Negative', 'Positive'], yticklabels=['Negative', 'Positive'])
plt.title('Confusion Matrix - Bidirectional GRU')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()

```

313/313 — 3s 9ms/step

Classification Report:

	precision	recall	f1-score	support
0	0.8943	0.8968	0.8955	5000
1	0.8965	0.8940	0.8953	5000
accuracy			0.8954	10000
macro avg	0.8954	0.8954	0.8954	10000
weighted avg	0.8954	0.8954	0.8954	10000



## ✓ Bidirectional LSTM

```









model_bi_lstm = Sequential([
    Embedding(input_dim=10000, output_dim=32, input_length=300),
    Bidirectional(LSTM(units=32, return_sequences=False)),
    Dropout(0.5),
    Dense(32),
    LeakyReLU(alpha=0.01),
    Dense(1, activation='sigmoid')
])

optimizer = Adam(learning_rate=0.00293)

model_bi_lstm.compile(
    loss='binary_crossentropy',
    optimizer=optimizer,
    metrics=['accuracy']
)

history_bi_lstm = model_bi_lstm.fit(
    X_train_pad, y_train,
    validation_data=(X_test_pad, y_test),
    epochs=60,
    batch_size=64,
    callbacks=[EarlyStopping(patience=5, restore_best_weights=True)],
    verbose=1
)

```

↩ Epoch 1/60  
 /usr/local/lib/python3.11/dist-packages/keras/src/layers/core/embedding.py:  
 warnings.warn(  
 /usr/local/lib/python3.11/dist-packages/keras/src/layers/activations/leaky\_  
 warnings.warn(  
**625/625**  **16s** 22ms/step - accuracy: 0.6776 - loss: 0.584  
 Epoch 2/60  
**625/625**  **13s** 21ms/step - accuracy: 0.8746 - loss: 0.324  
 Epoch 3/60  
**625/625**  **13s** 21ms/step - accuracy: 0.8830 - loss: 0.309  
 Epoch 4/60  
**625/625**  **13s** 21ms/step - accuracy: 0.9160 - loss: 0.228  
 Epoch 5/60  
**625/625**  **13s** 21ms/step - accuracy: 0.8935 - loss: 0.269  
 Epoch 6/60  
**625/625**  **13s** 21ms/step - accuracy: 0.9229 - loss: 0.208  
 Epoch 7/60  
**625/625**  **13s** 21ms/step - accuracy: 0.9425 - loss: 0.168  
 Epoch 8/60  
**625/625**  **13s** 22ms/step - accuracy: 0.9503 - loss: 0.145

```
y_pred_prob = model_bi_lstm.predict(X_test_pad).flatten()
y_pred = (y_pred_prob > 0.5).astype("int32")
```

↻ 313/313 ————— 3s 8ms/step


```
print(f" Accuracy: {accuracy_score(y_test, y_pred):.4f}")
print(f" Precision: {precision_score(y_test, y_pred):.4f}")
print(f" Recall:    {recall_score(y_test, y_pred):.4f}")
print(f" F1 Score:  {f1_score(y_test, y_pred):.4f}")
```

↻ Accuracy: 0.8822  
Precision: 0.8612  
Recall: 0.9112  
F1 Score: 0.8855

```

cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Oranges', xticklabels=['Negative', 'Positive'], yticklabels=['Negative', 'Positive'])
plt.title('Confusion Matrix - Bidirectional LSTM')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()

```



```

def plot_accuracy_loss(history, model_name="Model"):
    plt.figure(figsize=(14, 5))

    plt.subplot(1, 2, 1)
    plt.plot(history.history['accuracy'], label='Train Accuracy')
    plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
    plt.title(f'{model_name} - Accuracy Over Epochs')
    plt.xlabel('Epoch')
    plt.ylabel('Accuracy')
    plt.legend()


```



```
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title(f'{model_name} - Loss Over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

plt.tight_layout()
plt.show()
```

```
plot_accuracy_loss(history_bi_lstm, model_name="Bidirectional LSTM")
```



Start coding or [generate](#) with AI.

